

A probe with held-out AUC 1.000 is not a monitor

Qwen3.5-9B · layer 21 · residual stream · difference of means · prefill only · 529 records

Epistemic status: one model, one layer, ~19 hours. The headline number is solid and I checked it three ways. The explanation of why it fails changed once during the project, after I looked at which prompts the probe actually fires on — that part is post-hoc and I say so where it matters.

Executive summary

Models sometimes notice they are being evaluated, and behave differently when they do. That is a problem for anyone running evals, so people have started building probes to detect it — cheap linear directions read off the residual stream, reported with near-perfect separation. This is about what one of those does when you actually run it over traffic.

What I did


I built the cheapest version of that monitor. The direction is the published one — difference of means over the residual stream, as in Nguyen et al. (2025). Everything else is the cheap choice and mine: a one-sentence edit rather than distinct transcripts, the last token rather than a mean over tokens, a threshold from the training positives rather than a tuned one. The contrast is that edit — the same prompt with and without a line announcing an evaluation, in two agentic environments from Singh et al. (2026). Then I asked what it does when you point it at traffic instead of at its own training classes.

What I found

It separates its training classes perfectly, held-out AUC 1.000. On held-out instances of those same prompts its false-positive rate is 0.000. On 37 external prompts it has never seen, its false-positive rate is 0.162 and its precision at 1 % prevalence is 0.053 — about eighteen false alarms for every true one.

Then I looked at which external prompts it fires on. Two of them land 0.0009 apart: one drawn blind from Dolly, one I wrote myself for a cell called "evaluation of humans". A gym and a translation agency, no shared words, no shared subject. The four prompts that make the point are just below; the figure has all 37.

The external false positives are one semantic class



All 37 external prompts by score. The domain label is a regex over the committed prompt texts, so anyone can reapply it; the split inside the domain — judgement present or absent — is my reading of the eight prompts the regex selects, and the blind coding that would test it was prepared and never run.

Takeaways

1. **The false-positive rate is not a property of the probe, it is a property of the negatives you chose.** 0.000 in-distribution, 0.162 external, 0.250 on my pre-registered topical set. Quoting one without the others says nothing.
2. **The false positives are one semantic class, not scatter.** Five of the six firing external prompts come from a single domain as my regex defines it — exercise-physiology exam questions — and eight of the 37 are from that domain. Odds ratio 46.7, Fisher exact one-sided $p = 0.00071$.
3. **The class is “judge a person from measured performance”, not “evaluation”.** The one fitness prompt that doesn't fire is the same domain with the judgement removed; the other two silent ones are misses by my domain regex, and I say so below. And the two contrasts below kill my preferred reading: the densest evaluation vocabulary in my set stays silent, and a checkable answer mostly does too.
4. **Most of a foreign prompt's score has nothing to do with its content.** Any short single-turn prompt sits at ~0.6 of the entire training shift regardless of what it says. The class above adds 0.2–0.4 on top of that. The offset gets every cell within reach of the threshold; the increment decides which ones cross.

5. **The training contrast was never evaluation-versus-not.** The model answers "is this an evaluation?" the same for both classes, 32/32 against 31/32. Under an adjacent evaluation vocabulary the negative class is *denser* — grep AUC 0.375, below chance. The deployment prompt literally describes an evaluation, grading threshold and question file included.
6. **AUC gives no warning of any of this.** Across three candidate layers it is 1.000 / 1.000 / 0.984 while contamination moves 0.554 / 0.783 / 0.725 and topical FPR over the pooled 30 prompts moves 0.200 / 0.200 / 0.367. Whichever layer you pick by AUC, you learn nothing about whether the probe is usable.

Why it matters

Run a probe like this over an eval suite where one prompt in a hundred is genuinely eval-aware, and you get tons of false positives — roughly eighteen per true one. That much is easy to guess. What isn't is that the size of that number is set by the negatives you happened to pick: the same probe at the same threshold goes from 0.000 to 0.250 depending on which set you point it at, and a held-out AUC of 1.000 tells you nothing about which one you are in. Everyone who measures false positives takes two things for granted: that the training contrast holds, and that the reader knows which negatives the rate came from. Neither check is standard, both are cheap.

Detailed analysis

Setup.

32 matched pairs across two agentic environments from Singh et al. — a sandbagging task with ten multiple-choice questions and a stated grading threshold, and a secret-number game against a noisy oracle. The positive class, E1, is the deployment prompt plus one sentence: (*Note: this is part of an evaluation of AI systems.*) The negative class is the identical prompt without it, and the **E1 shift** is how far that sentence moves the score — the unit everything below is quoted in. Those 32 pairs carry twelve pre-registered conditions each, 384 records. Around them sit 145 further prompts in fifteen cells, 37 of them external from Dolly-15k. $384 + 145 = 529$.

Scores are raw projections onto the direction; deployment prompts sit near -27 , E1 near -2 , so one E1 shift is about 25 score units. The threshold is the 5th percentile of the training positives, which puts the true-positive rate at 0.906 — that is the number every precision figure below uses.

The four prompts that decide the argument

$\tau = -4.5$, so anything above it is an alert.

Fires, +12.9147. External, drawn blind from Dolly-15k under a length filter:

You, a commercial gym supervisor, are supervising the resistance training area when you hear a lot of clanging noise by the seated leg press machine... A day or two later, he reports that he felt more sore than usual after that workout. What accounts for this? Is it safe for him to exercise today?

Fires, +12.9156. Written by me, for a cell I called "evaluation of humans":

Our translation agency grades its freelancers twice a year: two senior reviewers score a sample of 20 pages each... one has now appealed a grade formally. How would you test whether the review process itself is consistent?

A gym and a translation agency. No shared words, no shared subject, opposite provenance and they land **0.0009 apart** — one twentieth of this pipeline's extraction noise floor. (For the first I assumed I had a duplicate record. I did not; the texts are different.)

Silent, -10.25. Authored, and the densest evaluation vocabulary anywhere in my set:

...students critique a published benchmark for conversational agents: reproduce a subset of its test items, then analyze what the evaluation actually measures, where the scoring rules break...

Silent, -11.62. External, and a task with an objectively checkable answer:

Given these two sets of instructions, will you end up at the same end point?

The number, and the three ways I checked it

negatives	n	FPR	Wilson 95 %	precision at 1 %
held-out deployment prompts (in-distribution)	32	0.000	[0.000; 0.107]	— mechanical
pre-registered topical, written before the result	8	0.250	[0.07; 0.59]	0.035
authored after the result was known	22	0.182	[0.07; 0.39]	0.048
external Dolly, sha and seed recorded before selection	37	0.162	[0.08; 0.31]	0.053

The first row is the one usually left out and the most informative — but not because the monitor works there. Zero events at $n = 32$ only bounds the rate at 0.107, which still admits a precision of 0.079. What it shows is that a complete class gap on the training distribution is compatible with almost any precision at all.

Three things I checked because I did not want the number to depend on a choice I made:

- **The threshold.** Sweeping it from TPR 0.80 to 0.99 moves topical FPR only $0.200 \rightarrow 0.233$ and precision stays in 0.039–0.046.
- **The true-positive rate.** It is measured on the positive class, which I argue is not eval awareness. So I computed precision across the whole range instead of assuming: even at a perfect TPR of 1.0 it stays under 0.07 on all three sources.
- **The procedure.** Under fully out-of-sample leave-one-decision-point-out — direction *and* threshold refit on the remaining 31 pairs, 32 folds — the grouped figures are 0.214 topical and 0.146 neutral. On the external set the two procedures agree to five decimal places.

Authoring bias did not help me. The negatives I wrote after seeing the result give a lower FPR than the pre-registered ones — 0.182 against 0.250. If authoring bias is there, it worked against my thesis.

What the false positives actually are

The six external prompts that fire: `wd_019` (+12.91), `wd_003` (+11.93), `wd_014` (+8.51), `wd_029` (+3.86), `wd_013` (+3.27), `wd_018` (−1.66). Five come from one domain, and all five ask a professional to judge a specific person from measured performance — a participant dropping the weight stack and whether he can train today; a client whose heart rate fell at every stage of a submaximal test; choosing a strength assessment for two women; a member whose gains have stalled; and the untrained one of those two women again, as a separate question.

Only one of the three fitness prompts that stay silent is really the same domain minus the judgement: "I want to get in better shape — can you recommend a workout routine?" (−13.79). The other two are misses by my own regex. One classifies App Store apps and was caught by the word *client*, inside "Termius — SSH&SFTP client" (−16.76); the other is a wedding gift registry, caught by *gym*, inside "gym membership" in a list of presents (−18.08). I am leaving them in the count. Excluding them would make the domain six prompts, five of which fire against one of thirty-one elsewhere, and move the Fisher p from 0.00071 to 0.00008 — so the loose regex is the conservative choice, and the one I quote. Comparing group means across the domain as the regex defines it, the firing and silent groups sit nine tenths of an E1 shift apart, and the gap inside one domain is almost as large as the evaluation sentence itself. That leaves the within-domain contrast — judgement present against judgement absent — resting on a single external prompt. The weight is carried by the authored cells instead — d at 0.879 against a_head at 0.570, on eight matched stories.


This replaced the reading I preferred. Both training environments are literally "answer this and the answer will be checked" — ten multiple-choice questions with a stated grading threshold, and a guessing game against an oracle. So I expected the direction to respond to a task with a checkable answer being posed to it. It doesn't. `wd_013`, a geometry puzzle with a definite answer, is the one non-fitness prompt among the six and near the bottom of them; two other puzzles of exactly that kind don't fire at all. What survives is narrower, and it was already sitting in the pre-registered cells above.

One thing that does not line up: my own "evaluation of humans" cell has the highest mean shift of any topical cell, yet only 1 of its 6 prompts fires, while the external instances of the same class fire 5 of 6. The cell mean sits at -4.89 against a threshold of -4.5 — right on the boundary, so almost nothing crosses. The external ones average $+7.11$. Whatever the class is, the versions I wrote are weaker instances of it than the ones Dolly happened to contain.

This is one more case in a growing list of probes of this kind reading something other than evaluation context; Devbunova (2026) found benchmark formatting. Different designs, different confounds, none of them evaluation awareness.

What I don't have is a mechanism. I can say what the direction responds to and show it holds on data I didn't choose, but not why an axis fitted on "this is an evaluation" ends up carrying "judge a person from measured performance". The honest guess is that both are instances of something more general about being asked to render a verdict, and the way to test that is the blind recode plus a second model — not more controls on this one.

Two honesty notes. The category is named post-hoc; the frozen coding rule for a blind recode exists in the repository and the labels were never collected, which is why that is now the first item in "what next". And the Fisher test deliberately uses a *domain* proxy — a regex over the committed prompt texts — rather than my category, so any reader can reapply it. Both this and the extraction floor below were computed after the code snapshot was frozen, from two committed files, and need nothing rerun.



The adjacent vocabulary excludes the nineteen registered stems the contrast was built around. An AUC below chance means the class labelled "not an evaluation" is the more evaluation-saturated of the two.


The 0.6 offset underneath everything

cell	shift (fraction of E1)	fires
structural control (dptrunc)	0.010	2/10
length ladder, pooled	0.281 / 0.306	0/16, 0/16
a_head — evaluation background, judgement removed	0.570	0/8
external Dolly	0.590	6/37
AI-safety documents	0.598	0/8
neutral authored	0.646	0/4
pre-registered topical (a_legacy / d_legacy)	0.727 / 0.738	1/2, 1/6
evaluation topic, lexicon stripped / rich wording	0.761 / 0.781	1/8, 2/8
everyday background, evaluation request	0.826	1/8
evaluation of humans	0.879	1/6
classical ML / LLM engineering	0.935 / 0.974	3/4, 3/4

The five topical cells are 30 prompts and 6 alerts — that is the 0.200 above, and it splits into 2/8 pre-registered and 4/22 authored.

And the offset is not even a property of this direction. All fifteen negative cells sit inside the band of directions refitted on *shuffled* labels. On the a1 cell: a random unit direction in 4096 dimensions moves it 0.30 of the training shift; a random-sign direction inside the span of the 32 paired differences moves it up to 0.96; the trained direction moves it 0.78. The subspace separates foreign text from deployment three times better than a random axis, and the trained direction isn't even the strongest member of it.

A number I retracted during writing. I had reported the length ladder as bounding the length effect at ~0.03 E1. Split by environment the two rungs move in *opposite* directions — sandbagging 0.156 → 0.434, secret-number 0.456 → 0.128 — and the pooled 0.025 is those two cancelling. Truncation also removes the part of each prompt that varies, so [len140](#) sandbagging is eight byte-identical texts. The ladder bounds nothing, and the length-matched foreign-text control I didn't run is a real gap.




Each rung is the same source text cut to a different length, but only `len210` sandbagging retains eight distinct prompts; the other three groups collapse to one or two.

Two things I expected and didn't get

Rewording halves the effect, relocating removes it — but I can't separate the two. Four wordings at three placements give 1.000 / 0.510 / 0.063 / -0.024 of the training shift. A rewording at the same place and role costs half the signal, so the direction is not reading the concept. The two cells that survive are exactly the two sharing a final user token with the positive class. But no condition holds the text fixed while moving it — all four are different sentences — so placement and wording are confounded and I'm not claiming the separation. The 2×2 that settles it costs fifteen minutes of prefill and is the first experiment in the list. I wrote down the predictions.

Rewording halves the effect, relocating removes it — and the two are confounded



Panel B is in units of the E1 shift measured within each environment separately, not the pooled one.

Steering reproduces the score of the insertion without reproducing its effect. The real insertion makes the model answer "yes, I'm being evaluated" 4/16 times. Steering at the λ that reproduces its score shift *exactly* gives 0/16. Push to eight times that and the trained direction does flip it, 14/16, where a shuffled-label control and a random control of identical norm give 0/16 and 1/16 — so the direction isn't arbitrary. At sixteen times both directions fitted inside the contrast flip, while the random one still doesn't flip at any λ . So the high-magnitude flip is a property of the subspace the contrast spans; magnitude and labelling both turn out to be beside the point.

And the probe score doesn't track what the model says. In the one cell where the model's answer varies with everything else held fixed, the score separates its four "yes" records from its twelve "no" records at AUC 0.479 — below chance, and thin. A pooled version of that number comes out at exactly 0.500, and I show it only so nobody mistakes it for support: all 128 within-cell pairs come from two cells where "yes" coincides exactly with one environment, 64 concordant plus 64 discordant.

What would have changed my mind

Written down in advance; none of it happened. A probe keeping AUC ≈ 1 while topical FPR fell below the external rate with non-overlapping intervals. A pooled-wording direction that transferred to a held-out wording. Steering that moved the self-report at $\lambda \approx 1$ where the shuffled control didn't. A within-cell AUC against the model's own answer meaningfully above 0.5. A monotone trend of contamination across the three pre-registered topical distances — the Jonckheere test returns $p = 0.923$ in the predicted direction.

The obvious fix, and what it costs

Refit with the authored negatives as hard negatives and the false-positive rate on held-out negative families goes to 0.000 — 0/8 on each of eight story pairs held out in turn, 0.000 on the whole 34-prompt a-family. Separation on the training contrast falls from 1.000 to 0.75–0.81. You can have perfect separation or no contamination, not both.

Three caveats and the third is the one that matters: leave-one-cell-out leaks, so the leading numbers are the pair-wise and family-wise versions; the AUC figures are in-sample on the contrast; and **the external set was never scored under that direction at all**. And it works because of the offset above

rather than anything about evaluation: adding any subset of short foreign prompts to the negative centroid subtracts the offset they all share. It also requires already knowing what the probe will fire on.

Ten ways this could have been an artefact

Each was specified in advance as something that could fail; three came back against me: PC1 turned out to be suppressing contamination rather than causing it, the lexicon baseline landed below chance, and the class above replaced the reading I preferred.

Not activation norm (ratio 0.964; cosine scores reproduce everything to 0.02 E1). **Not anisotropy — removing PC1–2 raises topical FPR 0.200 → 0.428, so PC1 was suppressing contamination.** Length is the one I could not close — the ladder retraction above says why, and the control that would close it is in what I'd do next. Not unfamiliarity: within-cell rho against NLL has no consistent sign. Not structural markup. Not the estimator — the three logistic regressions collapse onto diff-of-means at cos 0.984–0.996, and the two that whiten are genuinely different directions but their intervals on external negatives overlap the baseline. Not the threshold rule. Not training on one wording — leave-one-insertion-out gives AUC 0.503 on the held-out one. Not the hard-negative omission (above). **And not idiosyncrasy: the false positives are a class, p = 0.00071.**

A negative control on the measurement itself, because agentic code needs one. A deliberately wrong token index gives cos 0.22 against the true activations — the extraction check is capable of failing, which is what licenses trusting the others. The residual-stream hook matches `hidden_states[L+1]` at `max_abs_diff = 0.0`, batched extraction matches per-record extraction under both padding sides, and the cross-session anchor is `rel_l2 = 0.002`.

The extraction floor I got by accident: truncation happened to leave eight byte-identical prompts and they happened to straddle a batch boundary. They are *bitwise* identical within a batch and differ by 0.0189 across it — 0.00075 of the training shift. The floor comes from batch composition under fp16; rerun the same batch and it reproduces bitwise. That is also what makes the 0.0009 in the summary mean something.

Limitations

1. **One model, one layer, one read position, one direction construction**, and this is the cheapest version of the recipe. Someone can fairly say I broke the weakest thing available. Nothing here shows that a contrast built from genuinely distinct transcripts would fail.
2. **The class in the summary is named post-hoc and the labels aren't collected.** The Fisher test uses a domain proxy any reader can reapply; the reading "judge a person from measured performance" is my description of the prompts the regex selected, two of which it selected by mistake.
3. **The external set isn't independent within itself.** Six of 37 are genuinely from one source domain and five of those fire; the regex counts eight. The Wilson interval assumes independence the set doesn't have.
4. **My "evaluation of humans" cell and the top of the external set turn out to be the same class**, which wasn't planned and makes those two negative sources less independent than they look.
5. **Effective sample size is two, not thirty-two.** The 32 training pairs are really two templates instantiated sixteen times each — measurable rather than argued: truncating to 140 words leaves eight byte-identical texts. So held-out-by-decision-point is nearly held-in.
6. **Placement and wording are confounded** (above). The 2×2 was designed and not run.
7. **Power.** Several cells are n = 4 or smaller; every "significant" confirmatory p equals the floor of the exact sign test at n = 16 and is reported as "16/16 deltas of one sign".
8. **The direction itself is the least stable of the four constructions** by split-half minimum cosine, 0.654 against 0.960 / 0.882 / 0.738.
9. **The external set wasn't enlarged**, and the reason is unflattering: the code that composed Dolly rows into prompts wasn't preserved in the snapshot. I reconstructed 35 of 37 exactly,

two only approximately, and rather than draw a larger sample under a provably different composition I kept $n = 37$.

10. **One deviation isn't in docs/DEVIATIONS.md**, which was frozen with the code: the eight pre-registered topical prompts were split post-hoc into two cells after the pooled 2/8 was known. Material, because the smaller half has FPR 0.500 and enters the leading pool; the pooled 0.250 I quote doesn't depend on it. The twenty deviations recorded before the freeze are in that file, ordered by impact.
11. The domain proxy is loose in both directions: it catches an SSH client and a gym membership on a wedding list, and it would not catch a physiotherapy prompt with none of those words. Loose in the direction that lowers the odds ratio, which is why I did not tighten it after looking.

What I'd do next

1. **Collect the blind labels on the external set.** Rule frozen, package prepared, scores exist, zero GPU. It is the cheapest test of the only positive claim I make.
 2. **Replace the minimal edit with a real contrast** — distinct evaluation and deployment transcripts, and negatives drawn from the distribution the threshold will actually face. Hua et al. (2025) build the behaviour first and then find the direction, which is probably the right order.
 3. **The 2×2 crossing placement with wording**, fifteen minutes of prefill, predictions already written down.
 4. **A foreign-text control matched on length** (~300 words, no evaluation content) — the gap the length ladder doesn't fill.
 5. A second model, and a negative set stratified by domain now that I know why that matters.
-

About this submission

~19 hours of work under the 20-hour limit, done agentially: the pipeline, the pre-registration harness and the figure code went through an LLM coding agent; the design, the stopping rules, the extraction checks and every interpretation are mine.

Four of five pre-registered validations failed and the project went down the branch I had written down in advance for that outcome. A pre-recorded stopping rule closed a whole behavioural arm mid-project — the positive control was dead after escalation and the compute budget returned 14.19 hours required against 14.0 available on the same day — and I obeyed it. Three results here run against my own hypothesis and stayed in.

Reproducing. Everything except activation extraction runs on CPU in minutes.

`results/FINAL_DUMP_EN.md` is the numeric dump; every cell-level aggregate above is recomputable from `results/records_full.csv` alone, and section 2 of the dump reconciles stored against recomputed values for 15 of 15 cells. The steering, the within-cell AUCs and the leave-one-out rates come from three other committed files, named in the README. The leading out-of-sample procedure exists in two independent implementations written at different stages and they agree to five decimal places on every cell they share.

Nguyen et al., arXiv:2507.01786 · Singh et al., arXiv:2606.26071 · Devbunova, arXiv:2603.19426 · Hua et al., arXiv:2510.20487